



www.gisspecialisten.nl | van Deventerlaan 41 | 3528 AG Utrecht | +31 (0)85 7 325 825

Spatial Databases (1)

PostgreSQL+ PostGIS

Parya Pasha Zadeh
Utrecht, February 2020

Introduction and structure

Beginner	Introduction to databases and its advantages, Principles of relational data model, installation and data import, Database access through pgAdmin and QGIS Simple (spatial) SQL queries Views and tables
Intermediate	OGC simple features standards, Spatial data ingestion, Advanced spatial queries and analysis, Raster data handling in Postgres
Advanced	Integrity constraints implementation, Database functions and triggers, Server-side programming

Today's schedule

09:00 – 10:30	intro lecture relational databases, software installations materials access
10:30 – 10:45	<i>Coffee break</i>
10:45 – 12:30	non-spatial queries and SQL recap spatial data handling (spatial_ref_sys and geometry columns, OGC SFS)
12:30 – 13:15	<i>Lunch break</i>
13:15 – 15:00	spatial queries spatial functions invalid geometries
15:00 – 15:15	<i>Coffee break</i>
15:15 – 17:00	dynamic queries (views vs tables) PK-FK, GiST indexes The way forward

Objectives

Upon completion of this training you will be able to:

- Explain the concepts and structure of relational (spatial) databases;
- Understand the client-server architecture
- Recognize the differences between spatial and non-spatial DBMS
- Perform systematic query formulation using SQL
- Extend your SQL knowledge and use spatial functions in your queries.

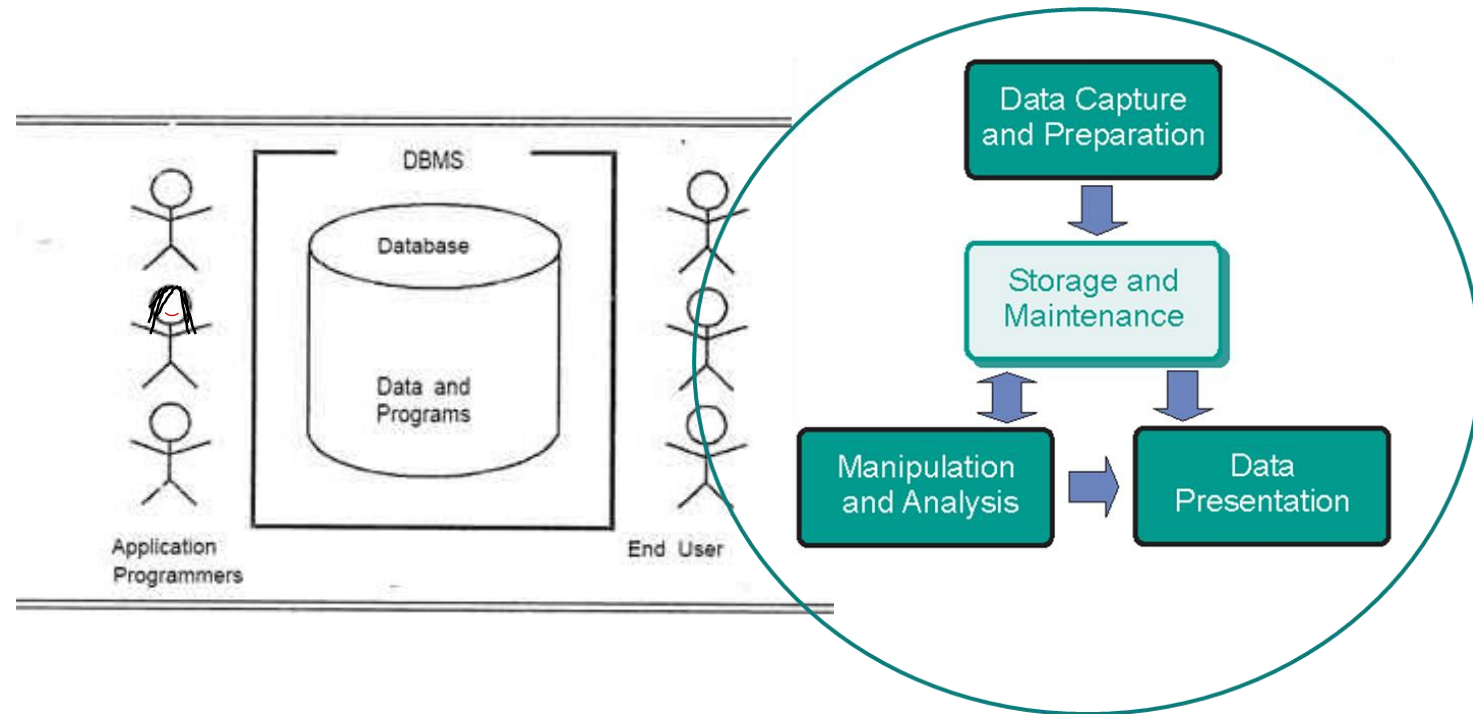
Data, Dataset, Database, DBMS and DB system

- Data
 - Is a resource held on paper or in digital format that serves to record or administer some facts and description of phenomena of interest.
- Data set (or dataset):
 - A homogeneous collection of data normally describing a single kind of phenomenon
- Database
 - *A collection of interrelated data sets properly structured* by means of, and stored through a DBMS.

Data, Dataset, Database, DBMS and DB system

- Database management system (DBMS)

- A *software package* that is designed for the purpose of managing databases. This means, DBMS allows to **set-up**, **maintain** and **interact with** one or more databases.
- DBMS uses a *data model* that is a collection of **data structuring primitives**, **rules of how to structure**, and **mechanisms to handle** the data in a database.



Database management system (DBMS)

- Supports the storage and manipulation of very large data sets.
- Can be instructed to guard over data correctness.
- Supports the concurrent use of the same data set by many users.
- Provides a high-level, declarative query language.
- Includes data backup and recovery functions to ensure data availability at all times.
- Allows the control of data redundancy.
- Supports the use of data model.




Data Model

A DBMS supports the use of a *data model*.


- Data model is an integrated collection of:
 - Data structuring primitives,
 - Rules of how to structure, and
 - Mechanisms to handle the data in a database.

In other words, a data model is a *toolbox* that allows us to *create/define* a database structure and *manipulate* the data stored in it.



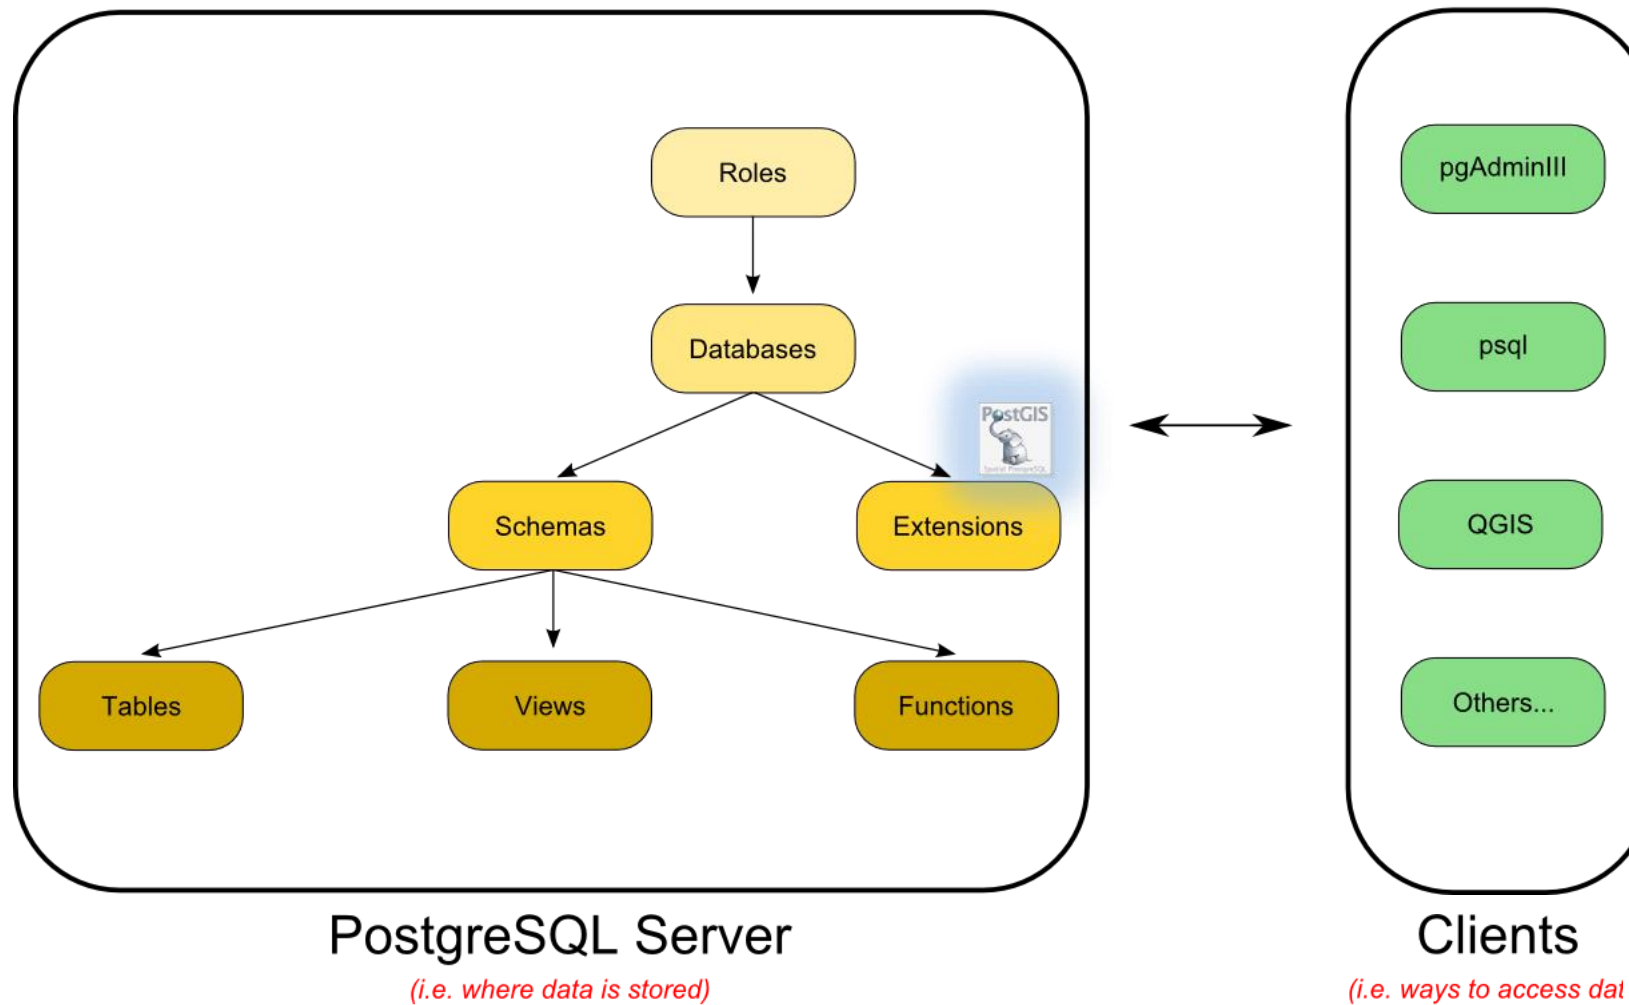
PostgreSQL and PostGIS

- **PostgreSQL** is a database management system
- It has a Server-Client architecture
- It runs on every main operating system
- It ensures purely **ACID** transactions:
 - **Atomicity** - (a transaction only knows two states – unstarted or finished)
 - **Consistency** - (a transaction always complies with the constraints of the DB)
 - **Isolation** - (how concurrent operations become visible to each other)
 - **Durability** - (after a transaction is terminated its effects remain)




PostgreSQL

PostgreSQL and PostGIS - System Architecture



PostgreSQL and PostGIS

- A spatial database allows users to store, query and manipulate collections of ***spatial data***.
- Spatial data is stored in as a geometry data type which complies with OGC standards.
- Spatial database is not constrained by the need to present data visually
- PostGIS is an extension that enable PostgreSQL with spatial data handling





Let's get started!

Materials and software

- Access the course at Github

https://github.com/paryapasha/Cursus_PostGIS_n1


- Open Source software

- PostgreSQL + PostGIS + pgAdmin

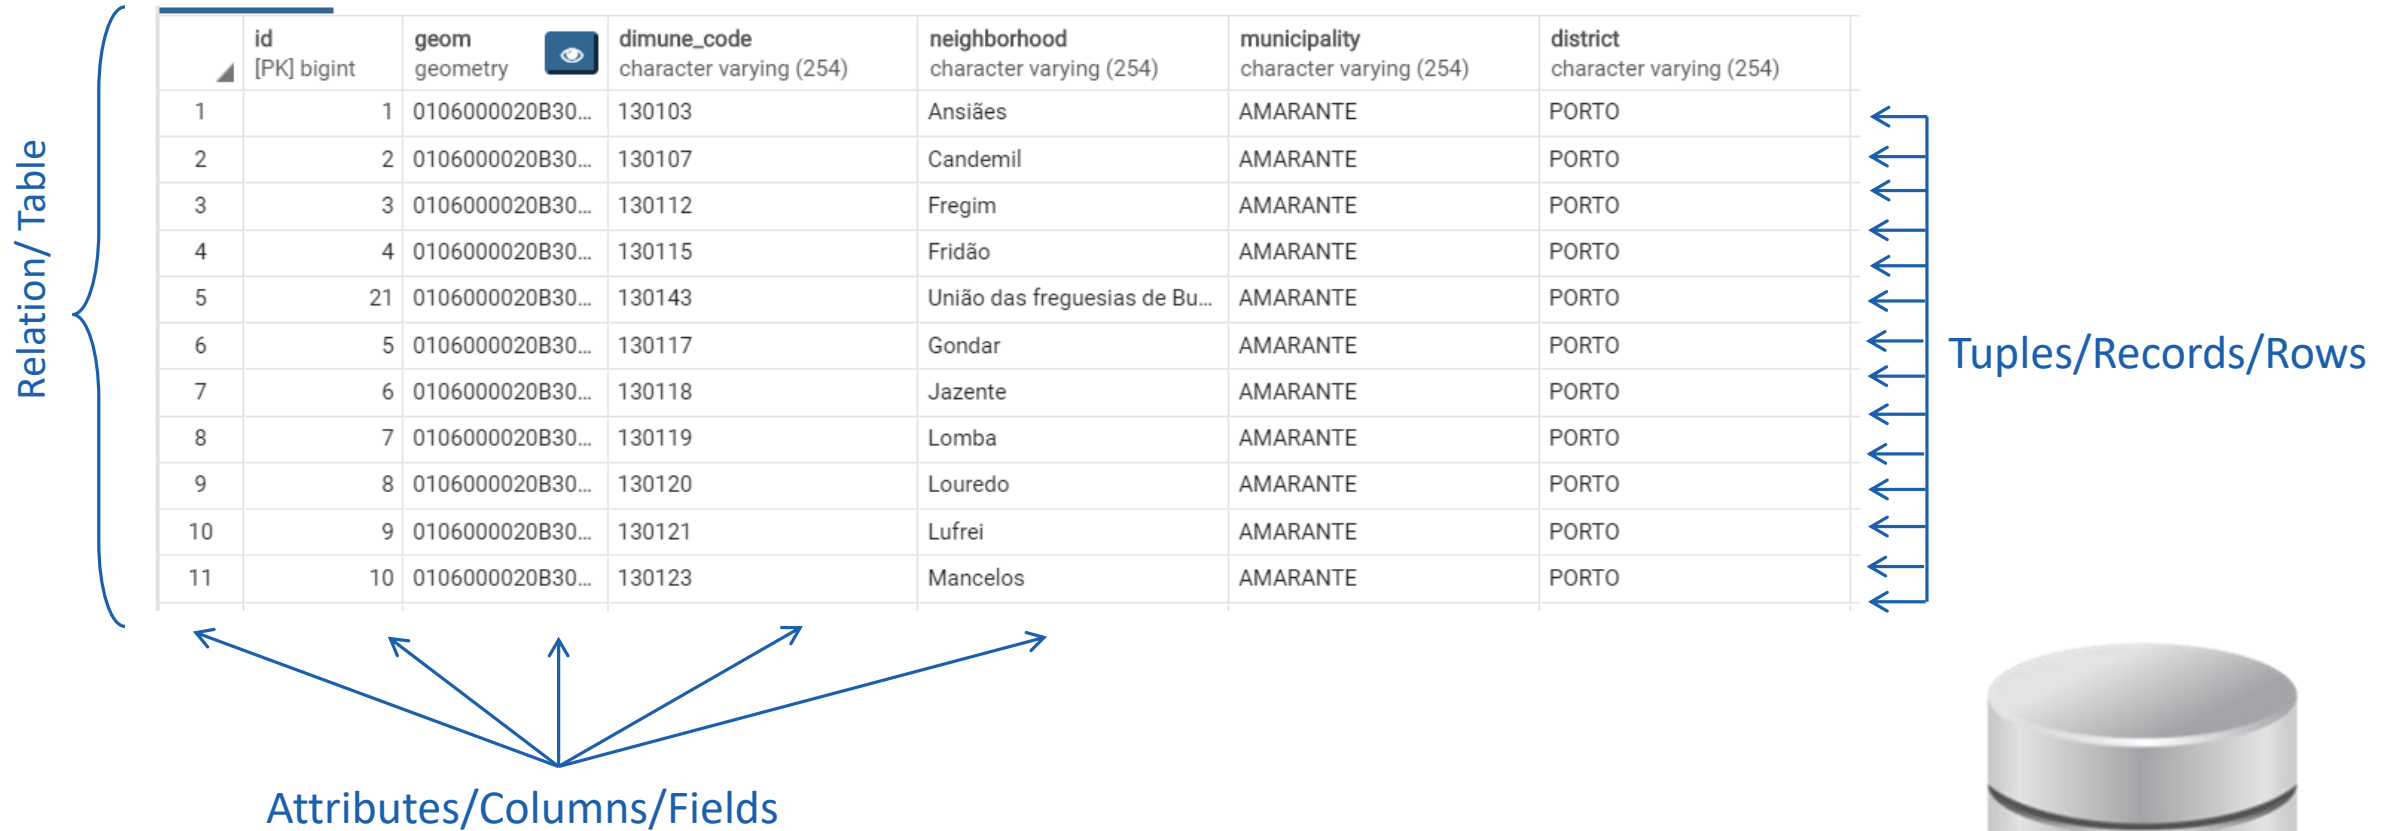
<https://www.enterprisedb.com/downloads/postgres-postgresql-downloads>

- QGIS (latest long term release 3.10)

<https://qgis.org/en/site/forusers/download.html#>



Query formulation



Query formulation

Simple Method of query definition

Steps:

1. INPUT RELATION

- On which relation is the query based?
- What will be the tuple variable?

2. SELECTION CONDITION

- What is the condition that the selected tuple for the output must fulfill?

3. OUTPUT RELATION

- What are the attributes in the output?

In SQL:

SELECT output tuple definition
(step 3)

FROM input tuple declaration
(step 1)

WHERE selection condition
(step 2)


JSP queries

Join, Select, Project

Very often, to extract required data from a database, there is a need to combine data from **two** or **more** relations in one query. In such query, selection condition depends on a **relationship of relations** involved in the query.

Phases of JSP query:

1. Cartesian product of input relations
 - *Which relations are involved in a query?*
2. Tuple selection
 - *What is the join condition?* ←
 - *What is the selection condition?*
3. Output tuple creation
 - *What attributes will the output relation have?*



Creating a table

Primary Key

- The key of a relation R is a set K of R's attributes such that:
 1. It is **unique** – there are no two distinct tuples of K that have the same attribute values.
 2. It is **minimal** – there is no proper subset of K that is unique.
 3. And it cannot be **NULL**

If there is more than one key for R, the database designer chooses one of them to become the primary key. The other remain candidate keys.



Referential integrity

Foreign Key

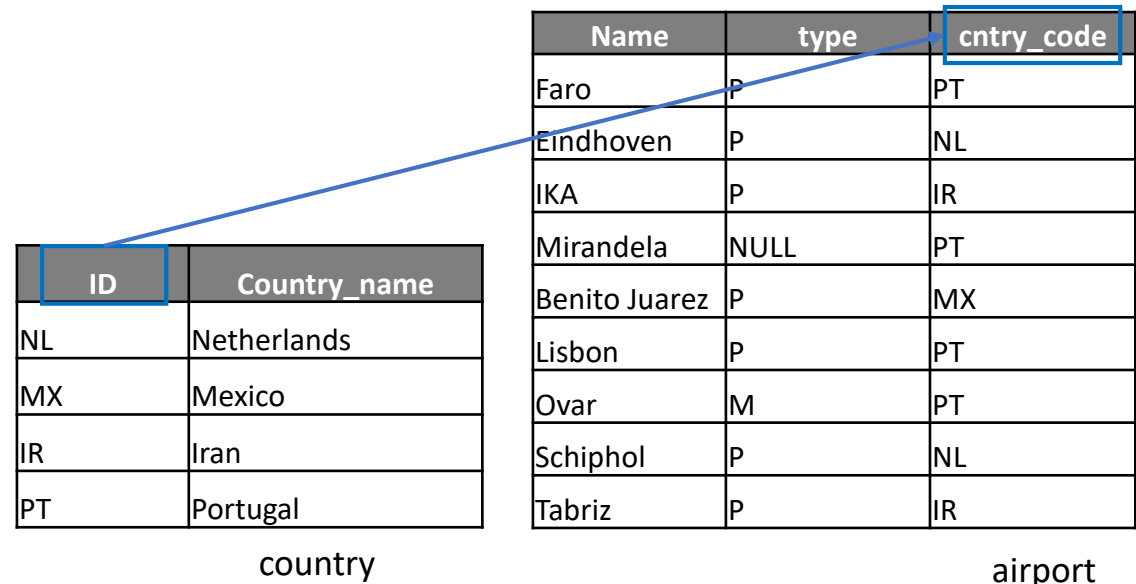
A foreign key is a set of attributes that is used to refer to a tuple in another relation.

- It must correspond with a primary key value in the second relation.
- A foreign key behaves like a 'logical pointer'.

```
SELECT *  
FROM table1, table2  
WHERE join condition
```

OR

```
SELECT *  
FROM table 1 INNER JOIN  
      table2 ON join condition
```



Referential integrity

Foreign Key

A foreign key is a set of attributes that is used to refer to a tuple in another relation.

- It must correspond with a primary key value in the second relation.
- A foreign key behaves like a 'logical pointer'.

```
SELECT a.type
FROM airport AS a, country AS c
WHERE c.ID = a.cntry_code
```

```
SELECT a.type
FROM airport AS a INNER JOIN
country AS c ON c.ID = a.cntry_code
```

ID	Country_name
NL	Netherlands
MX	Mexico
IR	Iran
PT	Portugal

Name	type	cntry_code
Faro	P	PT
Eindhoven	P	NL
IKA	P	IR
Mirandela	NULL	PT
Benito Juarez	P	MX
Lisbon	P	PT
Ovar	M	PT
Schiphol	P	NL
Tabriz	P	IR

country

airport

Summary queries

- A **summary query** is a query that infers statistical information from the tuples in the database.
- It involves **aggregate functions**: functions that operate on a set (bag) of tuples and produce a single value. Examples are:
 - counting tuples: *count*
 - summing up attribute values: *sum*
 - computing averages, minimum, maximum values from existing attribute values: *avg, min, max*
 - computing standard deviations, variances: *sdev, var*

```
SELECT attributename, aggregate (*)  
FROM tablename  
WHERE selection condition  
GROUP BY attributename  
HAVING groupselection condition
```



Spatial databases

Spatial data handling

What if you want to know...


- What is the area of a district?
- Do any of the countries in my table share a boundary?
- Which railroads pass through a city?
- How many schools are within a distance of 5 km to your house?

Spatially enabling Postgres

What happens when we enable PostGIS in our PostgreSQL database?

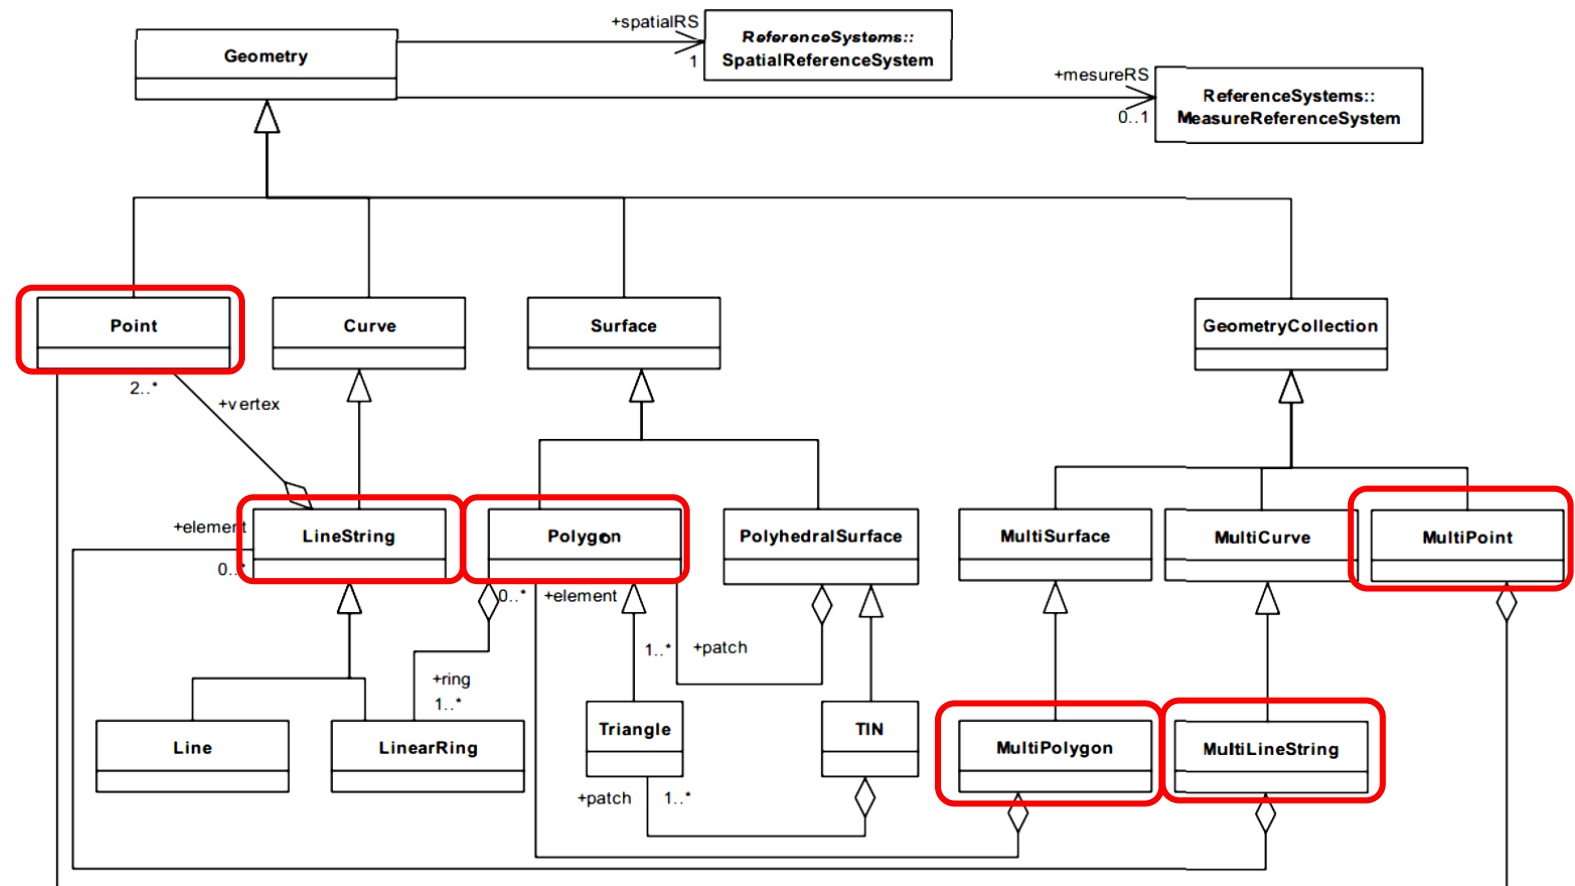
- Hundreds of *spatial functions* became available within the database;
- *Metadata views* are created:
 - geography columns
 - geometry columns
 - raster columns
 - raster overviews
- *Metadata table* is created:
 - spatial_ref_sys (SRIDs are the primary key to this table)

❖ All these metadata views and table are necessary in order to add support for OGC Simple Feature Model



OGC Simple Features Specification

- The Simple Feature Specification (SFS) is a standard from OGC on how data should be stored.
- It also defines functions for accessing, operating, and constructing data.




OGC Simple Features Specification

Geometry types


- Point

POINT (30 10)




- Linestring

LINESTRING (30 10, 10 30, 40 40)




- Polygon

POLYGON ((30 10, 40 40, 20 40, 10 20, 30 10))



- Polygon (with “hole”)

POLYGON ((35 10, 45 45, 15 40, 10 20, 35 10), (20 30, 35 35, 30 20, 20 30))




OGC Simple Features Specification

Geometry types


- MultiPoint

MULTIPOINT ((10 40), (40 30), (20 20), (30 10))




- MultiLinestring

MULTILINESTRING ((10 10, 20 20, 10 40), (40 40, 30 30, 40 20, 30 10))



- MultiPolygon

MULTIPOLYGON (((30 20, 45 40, 10 40, 30 20)), ((15 5, 40 10, 10 20, 5 10, 15 5)))




OGC Simple Features Specification

Geometry types


- MultiPoint

MULTIPOINT ((10 40), (40 30), (20 20), (30 10))




- MultiLinestring

MULTILINESTRING ((10 10, 20 20, 10 40), (40 40, 30 30, 40 20, 30 10))



- MultiPolygon

MULTIPOLYGON (((30 20, 45 40, 10 40, 30 20)), ((15 5, 40 10, 10 20, 5 10, 15 5)))



Some famous (!) Spatial functions

ST_Length

```
SELECT id, st_length(geom)
FROM rivers
WHERE rivername = 'IJssel'
```

output: length of the
line/multiline

ST_Area

```
SELECT st_area(s.geom)
FROM squares as s
```

output: area of the
polygon/multipolygon

ST_Centroid


```
SELECT b.id, st_centroid (b.geom)
FROM buildings as b
```

output: a point geometry,
centroid of input geometry

Some famous (!) Spatial functions


ST_Buffer

```
SELECT id, st_buffer(geom, 100)
FROM    rivers
WHERE    rivername = 'Tejo'
```



ST_Within (check st_dwithin and st_contains)

```
SELECT l.id, st_within(l.geom, s.geom)
FROM    lamps as l, squares as s
```



ST_Intersection (check st_intersects)

```
SELECT b.id, st_intersection(b.geom, s.geom)
FROM    buildings as b, squares as s
```



More Spatial function to explore

- `ST_transform`
- `ST_distance`
- `ST_dwithin`
- `ST_crosses`
- `ST_overlaps`
- `ST_disjoint`
- `ST_contains`
- ...

In addition to the functionality, try to understand the number and type of input features and the type of output.





On DB design and integrity

Creating a table

```
CREATE TABLE vectors.porto_neighborhood
(
    id bigint NOT NULL,
    geom geometry(MultiPolygon,3763),
    dimune_code character varying(254) ,
    neighborhood character varying(254) ,
    municipality character varying(254) ,
    district character varying(254) ,
    area_ha character varying(254) ,
    neighborhood_shortcode character varying(254) ,
    CONSTRAINT porto_neighborhood_pkey PRIMARY KEY (id)
)
```

Create Views

View vs. Table

- A view is a stored query in a form of a table which is not physically materialized;
- A view table can be queried but cannot be directly updated;
- A view changes every time the underlying data changes;
- Views are therefore dynamic;

```
CREATE VIEW expansion AS  
    SELECT id, area, st_buffer(geom, 10)  
    FROM landplot;
```



Spatial indexing

- An important feature of spatial databases is the support of Spatial Indexes.
- PostGIS uses GIST Indexation.

- Can simply be created on a table

CREATE INDEX index_name **ON** tablename

USING GIST (geomcolumnname);



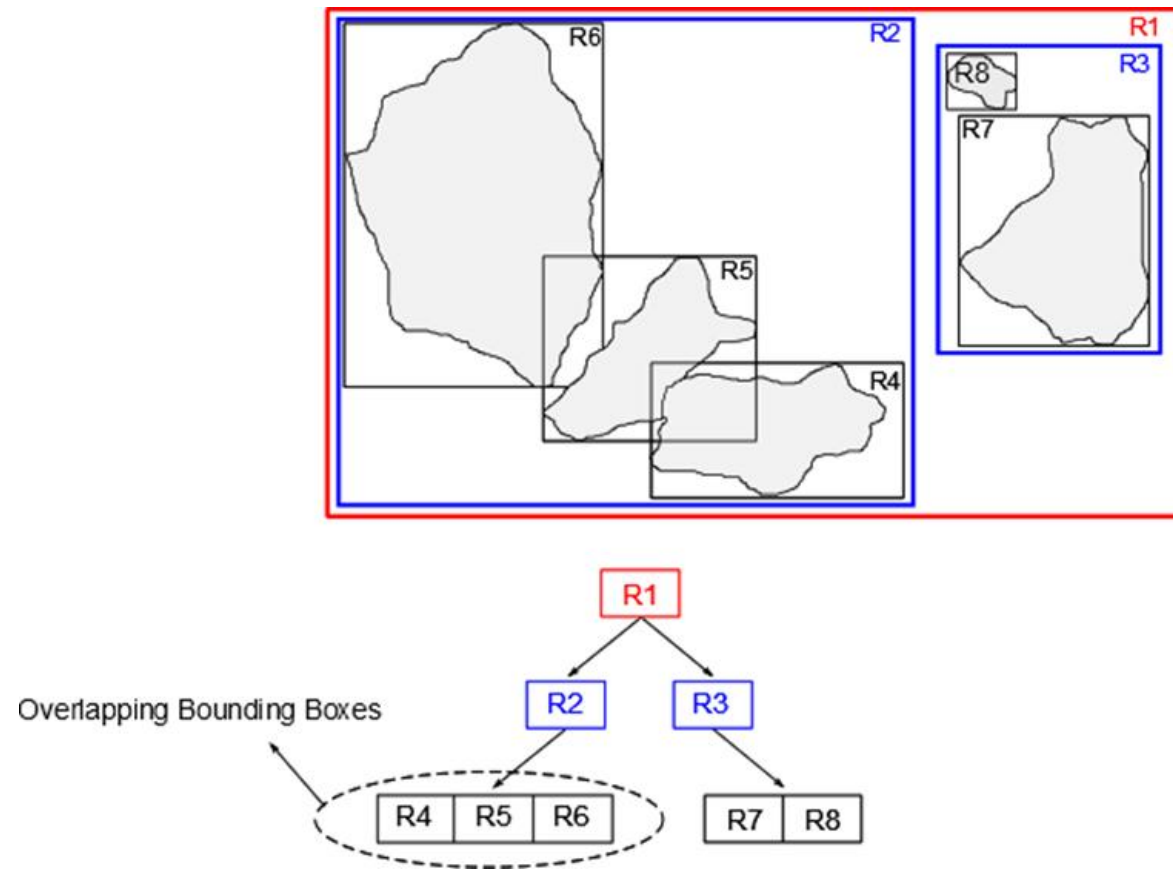
Spatial indexing

- Which features intersect (`st_intersects`)?



Spatial indexing

- First the database starts by identifying the overlapping bounding boxes:



Spatial indexing

- Then the actual topological operation takes place on the geometries whose bounding boxes overlap, giving us the result




- Most of the commonly used functions in PostGIS (ST_Contains, ST_Intersects, ST_DWithin, etc) include an index filter automatically. But some functions (e.g., ST_Relate) do not include an index filter.
- To do a bounding-box search using the index (and no filtering), make use of the **&&** operator.
- For geometries, the **&&** operator means “bounding boxes overlap or touch” in the same way that for numbers the = operator means “values are the same”.



How to continue...

Continue

- PostGIS in Action 2nd edition
 - Find out all about it [here](#)
 - Download the code and the data [here](#)
 - Purchase [here](#)
- Simple Feature Access
 - Part 1: [Common Architecture](#) (pg 25-40)
 - Part 2: [SQL Option](#)
- Manuals
 - [PostgreSQL](#)
 - [PostGIS](#)



Lesley Porter

Office Manager

M: +31(0)6 11 12 07 47

E: lesley@gisspecialisten.nl

Parya Pasha Zadeh

Operations Manager


M: +31(0)6 83 84 72 95

E: parya@gisspecialisten.nl

Spatial vs non-spatial queries




```
neighbourhood (  
  neigh_name String,  
  pop Double,  
  geom (Polygon,28992)  
)
```



```
bus_stop(  
  stop_name String,  
  geom (Point,28992)  
)
```

Spatial vs non-spatial queries



```
SELECT bus_stop.stop_name, neighborhoods.neigh_name,  
       neighborhoods.pop  
FROM neighborhoods JOIN bus_stop  
ON     ST_Intersects(neighborhoods.geom, bus_stop.geom)
```

```
SELECT *  
FROM table1, table2  
WHERE join condition
```

OR

```
SELECT *  
FROM table 1 INNER JOIN  
      table2 ON join condition
```